

Welcome To ...

Data Structures

王惠嘉



Sorting

- Rearrange n elements into ascending order.
- 7, 3, 6, 2, 1 \rightarrow 1, 2, 3, 6, 7

Insertion Sort



- $n \leq 1 \rightarrow$ already sorted. So, assume $n > 1$.
- $a[0:n-2]$ is sorted recursively.
- $a[n-1]$ is inserted into the sorted $a[0:n-2]$.
- Complexity is $O(n^2)$.
- Usually implemented nonrecursively (see text).

Quick Sort

- When $n \leq 1$, the list is sorted.
- When $n > 1$, select a **pivot** element from out of the n elements.
- Partition the n elements into 3 segments **left**, **middle** and **right**.
- The **middle** segment contains only the **pivot** element.
- All elements in the **left** segment are \leq **pivot**.
- All elements in the **right** segment are \geq **pivot**.
- Sort **left** and **right** segments recursively.
- Answer is sorted **left** segment, followed by **middle** segment followed by sorted **right** segment.

Example

6	2	8	5	11	10	4	1	9	7	3
---	---	---	---	----	----	---	---	---	---	---

Use 6 as the pivot.

2	5	4	1	3	6	7	9	10	11	8
---	---	---	---	---	---	---	---	----	----	---

Sort left and right segments recursively.

Choice Of Pivot


- Pivot is leftmost element in list that is to be sorted.
 - When sorting `a[6:20]`, use `a[6]` as the pivot.
 - Text implementation does this.
- Randomly select one of the elements to be sorted as the pivot.
 - When sorting `a[6:20]`, generate a random number `r` in the range `[6, 20]`. Use `a[r]` as the pivot.

Choice Of Pivot

- Median-of-Three rule. From the leftmost, middle, and rightmost elements of the list to be sorted, select the one with median key as the pivot.
 - When sorting $a[6:20]$, examine $a[6]$, $a[13]$ ($(6+20)/2$), and $a[20]$. Select the element with median (i.e., middle) key.
 - If $a[6].key = 30$, $a[13].key = 2$, and $a[20].key = 10$, $a[20]$ becomes the pivot.
 - If $a[6].key = 3$, $a[13].key = 2$, and $a[20].key = 10$, $a[6]$ becomes the pivot.

Choice Of Pivot

- If $a[6].key = 30$, $a[13].key = 25$, and $a[20].key = 10$, $a[13]$ becomes the pivot.
- When the pivot is picked at random or when the median-of-three rule is used, we can use the quick sort code of the text provided we first swap the leftmost element and the chosen pivot.



Partitioning Example Using Additional Array

a

6	2	8	5	11	10	4	1	9	7	3
---	---	---	---	----	----	---	---	---	---	---

b

2	5	4	1	3	6	7	9	10	11	8
---	---	---	---	---	---	---	---	----	----	---

Sort left and right segments recursively.

In-Place Partitioning Example

a

6	2	8	5	11	10	4	1	9	7	3
---	---	---	---	----	----	---	---	---	---	---

a

6	2	3	5	11	10	4	1	9	7	8
---	---	---	---	----	----	---	---	---	---	---

a

6	2	3	5	1	10	4	11	9	7	8
---	---	---	---	---	----	---	----	---	---	---

a

6	2	3	5	1	4	10	11	9	7	8
---	---	---	---	---	---	----	----	---	---	---

bigElement is not to left of **smallElement**,
terminate process. Swap **pivot** and **smallElement**.

a

4	2	3	5	1	6	10	11	9	7	8
---	---	---	---	---	---	----	----	---	---	---

Merge Sort

- Partition the $n > 1$ elements into two smaller instances.
- First $\text{ceil}(n/2)$ elements define one of the smaller instances; remaining $\text{floor}(n/2)$ elements define the second smaller instance.
- Each of the two smaller instances is sorted recursively.
- The sorted smaller instances are combined using a process called merge.
- Complexity is $O(n \log n)$.
- Usually implemented nonrecursively.

Merge Two Sorted Lists

- $A = (2, 5, 6)$
 $B = (1, 3, 8, 9, 10)$
 $C = ()$
- Compare smallest elements of A and B and merge smaller into C .
- $A = (2, 5, 6)$
 $B = (3, 8, 9, 10)$
 $C = (1)$


Merge Two Sorted Lists

- $A = (5, 6)$
 $B = (3, 8, 9, 10)$
 $C = (1, 2)$
- $A = (5, 6)$
 $B = (8, 9, 10)$
 $C = (1, 2, 3)$
- $A = (6)$
 $B = (8, 9, 10)$
 $C = (1, 2, 3, 5)$


Merge Two Sorted Lists

- $A = ()$
 $B = (8, 9, 10)$
 $C = (1, 2, 3, 5, 6)$
- When one of A and B becomes empty, append the other list to C .
- $O(1)$ time needed to move an element into C .
- Total time is $O(n + m)$, where n and m are, respectively, the number of elements initially in A and B .

Merge Sort



Merge Sort




Time Complexity

- Let $t(n)$ be the time required to sort n elements.
- $t(0) = t(1) = c$, where c is a constant.
- When $n > 1$,
$$t(n) = t(\text{ceil}(n/2)) + t(\text{floor}(n/2)) + dn,$$
where d is a constant.
- To solve the recurrence, assume n is a power of 2 and use repeated substitution.
- $t(n) = O(n \log n)$.

Nonrecursive Version

- Eliminate downward pass.
- Start with sorted lists of size 1 and do pairwise merging of these sorted lists as in the upward pass.

Nonrecursive Merge Sort



Complexity

- Sorted segment size is 1, 2, 4, 8, ...
- Number of merge passes is $\text{ceil}(\log_2 n)$.
- Each merge pass takes $O(n)$ time.
- Total time is $O(n \log n)$.
- Need $O(n)$ additional space for the merge.
- Merge sort is slower than insertion sort when $n \leq 15$ (approximately). So define a **small instance** to be an instance with $n \leq 15$.
- Sort small instances using insertion sort.
- Start with segment size = 15.

Natural Merge Sort

- Initial sorted segments are the naturally occurring sorted segments in the input.
- Input = [8, 9, 10, 2, 5, 7, 9, 11, 13, 15, 6, 12, 14].
- Initial segments are:
[8, 9, 10] [2, 5, 7, 9, 11, 13, 15] [6, 12, 14]
- 2 (instead of 4) merge passes suffice.
- Segment boundaries have $a[i] > a[i+1]$.

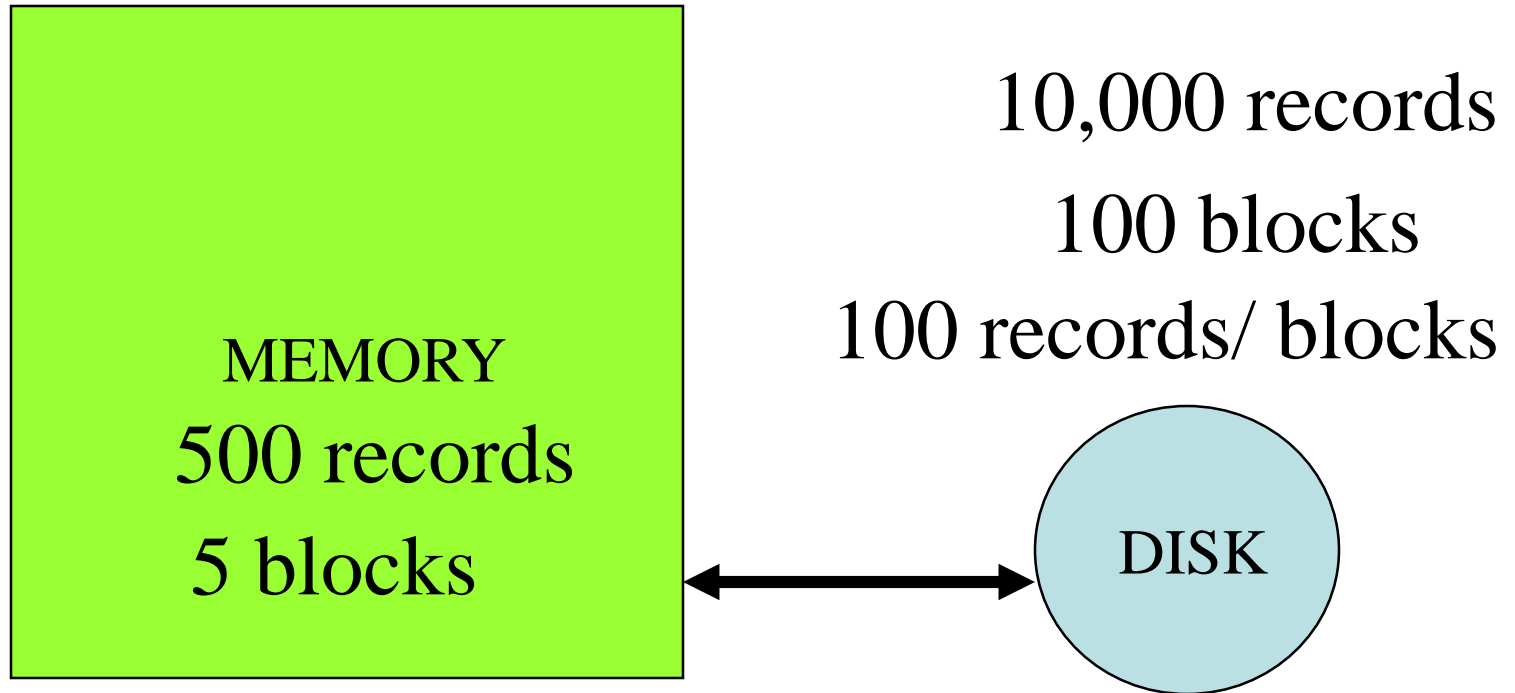
External Merge Sort

- Sort 10,000 records.
- Enough memory for 500 records.
- Block size is 100 records.
- t_{IO} = time to input/output 1 block
(includes seek, latency, and transmission times)
- t_{IS} = time to internally sort 1 memory load
- t_{IM} = time to internally merge 1 block load

External Merge Sort


- Two phases.
 - Run generation.
 - A run is a sorted sequence of records.
 - Run merging.

Run Generation



- Input 5 blocks.
- Sort.
- Output 5 blocks as a run.
- Do 20 times.

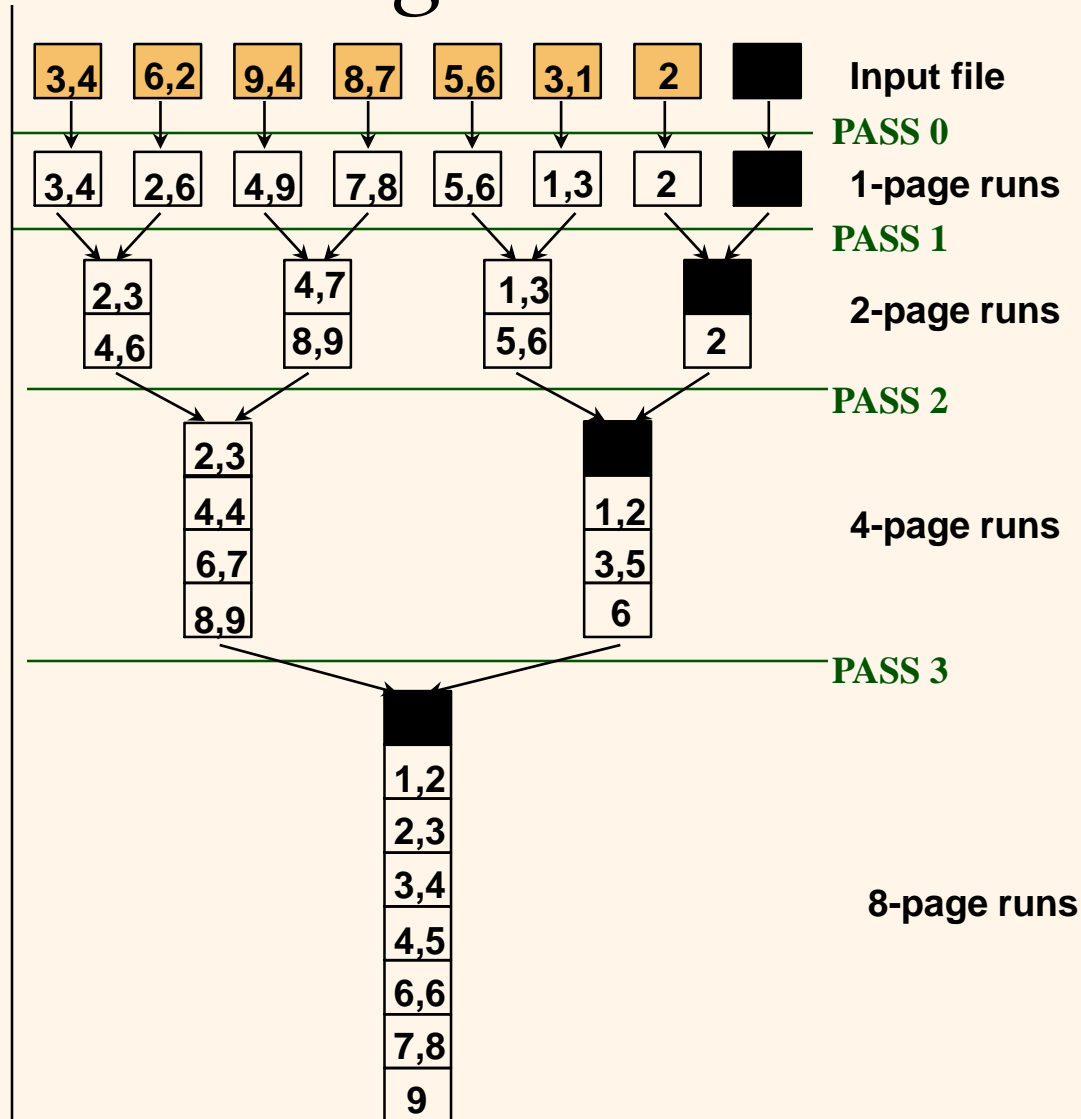
- $5t_{IO}$
- t_{IS}
- $5t_{IO}$
- $200t_{IO} + 20t_{IS}$



Two-Way External Merge Sort

- ❖ Each pass we read + write each page in file.
- ❖ N pages in the file \Rightarrow the number of passes $= \lceil \log_2 N \rceil + 1$
- ❖ So total cost is:


$$2N(\lceil \log_2 N \rceil + 1)$$
- ❖ Idea: *Divide and conquer*: sort subfiles and merge




Run Merging

- Merge Pass.
 - Pairwise merge the 20 runs into 10.
 - In a merge pass all runs (except possibly one) are pairwise merged.
- Perform 4 more merge passes, reducing the number of runs to 1.

Merge 20 Runs



Merge R1 and R2



- Fill l_0 (Input 0) from R_1 and l_1 from R_2 .
- Merge from l_0 and l_1 to output buffer.
- Write whenever output buffer full.
- Read whenever input buffer empty.

Time To Merge R1 and R2

- Each is 5 blocks long.
- Input time = $10t_{IO}$.
- Write/output time = $10t_{IO}$.
- Merge time = $10t_{IM}$.
- Total time = $20t_{IO} + 10t_{IM}$.

Time For Pass 1 ($R \rightarrow S$)

- Time to merge one pair of runs
 $= 20t_{IO} + 10t_{IM}$.
- Time to merge all 10 pairs of runs
 $= 200t_{IO} + 100t_{IM}$.

Time To Merge S1 and S2

- Each is 10 blocks long.
- Input time = $20t_{IO}$.
- Write/output time = $20t_{IO}$.
- Merge time = $20t_{IM}$.
- Total time = $40t_{IO} + 20t_{IM}$.

Time For Pass 2 (S \rightarrow T)

- Time to merge one pair of runs
 $= 40t_{IO} + 20t_{IM}$.
- Time to merge all 5 pairs of runs
 $= 200t_{IO} + 100t_{IM}$.

Time For One Merge Pass

- Time to input all blocks = $100t_{IO}$.
- Time to output all blocks = $100t_{IO}$.
- Time to merge all blocks = $100t_{IM}$.
- Total time for a merge pass = $200t_{IO} + 100t_{IM}$.

Total Run-Merging Time


- (time for one merge pass) * (number of passes)
= (time for one merge pass)
 * $\text{ceil}(\log_2(\text{number of initial runs}))$
= $(200t_{IO} + 100t_{IM}) * \text{ceil}(\log_2(20))$
= $(200t_{IO} + 100t_{IM}) * 5$

Factors In Overall Run Time

- Run generation. $200t_{IO} + 20t_{IS}$
 - Internal sort time.
 - Input and output time.
- Run merging. $(200t_{IO} + 100t_{IM}) * \text{ceil}(\log_2(20))$
 - Internal merge time.
 - Input and output time.
 - Number of initial runs.
 - Merge order (number of merge passes is determined by number of runs and merge order)

Improve Run Generation

- Overlap input, output, and internal sorting.




Improve Run Generation

- Generate runs whose length (on average) exceeds memory size.
- Equivalent to reducing number of runs generated.

Improve Run Merging

- Reduce number of merge passes.
 - Use higher-order merge.
 - Number of passes
= $\text{ceil}(\log_k(\text{number of initial runs}))$
where k is the merge order.

Merge 20 Runs Using 4-Way Merging



Number of merging passes = 3

Total passes = 1 (run generation) + 3 (merging) = 4₄₄

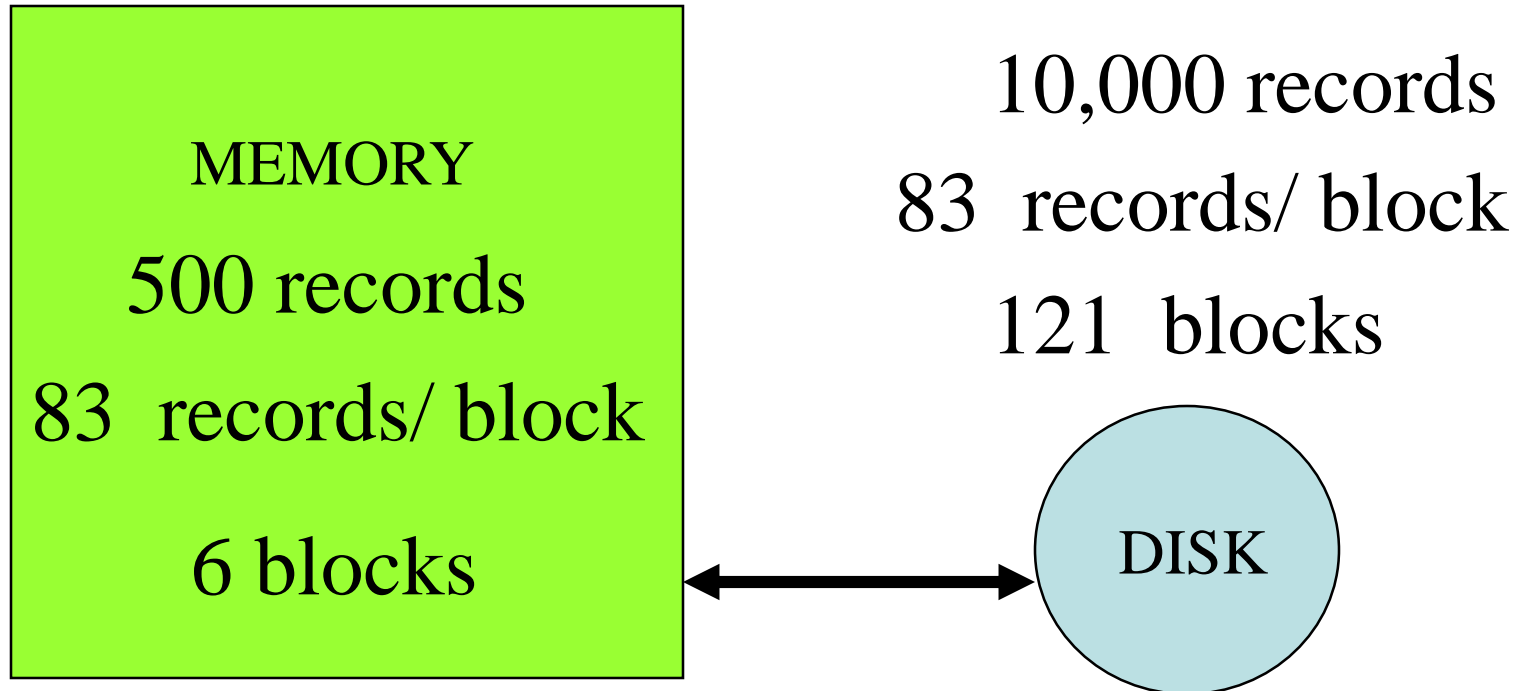
Time For Pass 1 (R S)

- Time to merge 4-way of runs
 $40t_{IO} + 20t_{IM}$.
- Time to merge all runs (5 blocks/run)
 $200t_{IO} + 100t_{IM}$.

Total Run-Merging Time

- (time for one merge pass) * (number of passes)
= (time for one merge pass)
* $\text{ceil}(\log_4(\text{number of initial runs}))$
= $(200t_{IO} + 100t_{IM}) * \text{ceil}(\log_4(20))$
= $(200t_{IO} + 100t_{IM}) * 3$


Run Generation



- Input 6 blocks.
- Sort.
- Output 6 blocks as a run.
- Do 21 times.

- $6 t_{IO}$
- t_{IS}
- $6 t_{IO}$
- $242t_{IO} + 21t_{IS}$

Merge 21 Runs Using 5-Way Merging



Number of merging passes = 2

Total passes = 1 (run generation) + 2 (merging)
= 3


Time For Pass 1 (R \rightarrow S)

- Time to merge 5-way of runs
 $60t_{IO} + 30t_{IM}$.
- Time to merge all runs (6 blocks/run)
 $242t_{IO} + 121t_{IM}$.

Total Run-Merging Time

- (time for one merge pass) * (number of passes)
= (time for one merge pass)
 * $\text{ceil}(\log_5(\text{number of initial runs}))$
= $(242t_{IO} + 121t_{IM}) * \text{ceil}(\log_5(21))$
= $(242t_{IO} + 121t_{IM}) * 2$


Run Generation



- Input **10** blocks.
- Sort.
- Output **10** blocks as a run.
- Do **20** times.

- $10 t_{IO}$
- t_{IS}
- $10 t_{IO}$
- $400t_{IO} + 20t_{IS}$

Merge 20 Runs Using 9-Way Merging



Number of merging passes = 2

Total passes = 1 (run generation) + 2 (merging)
= 3

Time For Pass 1 (R \rightarrow S)

- Time to merge 9-way of runs
 $180t_{IO} + 90t_{IM}$.
- Time to merge all runs
 $800t_{IO} + 400t_{IM}$.

Total Run-Merging Time


- (time for one merge pass) * (number of passes)
= (time for one merge pass)
* $\text{ceil}(\log_2(\text{number of initial runs}))$
= $(400t_{IO} + 200t_{IM}) * \text{ceil}(\log_2(20))$
= $(400t_{IO} + 200t_{IM}) * 2$

I/O Time Per Merge Pass

- Number of input buffers needed is linear in merge order k .
- Since memory size is fixed, block size decreases as k increases (after a certain k).
- So, number of blocks increases.
- So, number of seek and latency delays per pass increases.

I/O Time Per Merge Pass

I/O
time
per
pass




k

Total I/O Time To Merge Runs



- (I/O time for one merge pass)
* $\text{ceil}(\log_k(\text{number of initial runs}))$)

Total
I/O
time to
merge
runs




k

Internal Merge Time



- Naïve way $\Rightarrow k - 1$ compares to determine next record to move to the output buffer.
- Time to merge n records is $c(k - 1)n$, where c is a constant.
- Merge time per pass is $c(k - 1)n$.
- Total merge time is $c(k - 1)n \log_k r$.

Merge Time Using A Selection Tree



- Time to merge n records is $dn\log_2 k$, where d is a constant.
- Merge time per pass is $dn\log_2 k$.
- Total merge time is $(dn\log_2 k) \log_k r = dn\log_2 r$.

Improve Run Merging

- Reduce number of merge passes.
 - Use higher order merge.
 - Number of passes
= $\text{ceil}(\log_k(\text{number of initial runs}))$
where k is the merge order.
- More generally, a higher-order merge reduces the cost of the optimal merge tree.